

Filesystem Storage Deduplication

IN4120 — Search Technology

Storage

- Modern capacities are still filled up
- Different applications, different needs
- Redundant backups
- Different users, identical data

Our goal

- **Find** duplicate files on live system
- Find duplicate **directories**
- **Avoid** insertion of already-existing files
- **Delete** unnecessary copies

Why bother?

Performance hits



- **Indexing** is not free
- Taking **backups** takes time
- **Virus** detection takes CPU time
- Overhead during **transactions**
- ...

Economical costs

- Using **bandwidth**
- Upgrading storage
- ...

More reasons: Data Management

Systems perspective

- Identifying original version
 - *"Fuck, I deleted the wrong version"*
 - Overhead during identification
 - Both for **computers** 
 - and our silly **brains** 
- Validation

Data perspective

- Complexity during compression
- Entropy loss
- Two-phased (de)compression
- Copy-on-Write (CoW) latency

Benefits of free storage space

Operating system

- Swap space

File system

- Idle time tidying
 - Defragmentation
- Prefetching buffer

Identity comparison

- Early exit
- Hashing
 - Large files: preemptive comparison of fixed-size bytes
 - Small files: Bytewise?
- Metadata
 - Filesize
 - Filename?
- Reading from IO devices is SLOW
- Many-to-many comparisons

Comparing directories

Using hashes as content, recursively

- Destructive: loses meaningful information
 - Inability to do similarity comparisons

Metadata

- Number of files/directories
- Filetypes included

Top-bottom (breadth-first) vs Merkle tree

Dynamic systems

Bottom-up propagation approach

Merkle tree invalidation

- Invalidate parents on file change
- Avoid recomputing hashes
 - Blocking operation, adds write latency
 - May not be read before invalidated again

Filesystem support

- Hash information as xattr
- Computed during idle time
- Available from superblock metadata

```
~  
doas python depthcount.py /  
Depth      Files      Dirs  
  1           0        19  
  2         941       1_009  
  3        43_405      8_330  
  4       184_602     16_679  
  5       321_794     50_555  
  6    2_549_515     97_574  
  7       681_773     54_903  
  8       562_982     53_385  
  9       428_951     87_255  
 10       540_964     56_997  
 11       439_686     51_933  
 12       358_841     51_009  
 13       334_918     40_280  
 14       318_916     27_913  
 15       208_358     17_960  
 16       163_164     13_396  
 17        80_622     10_462  
 18        56_621      5_413  
 19        22_449      3_396  
 20        20_684      4_879  
 21        10_604      2_404  
 22         3_768      1_514  
 23         2_122       163  
 24          389        50  
 25           74        24  
 26           20         4  
 27            4         0
```

```
~  
python depthcount.py ~  
Depth      Files      Dirs  
  1         128        62  
  2         258       420  
  3        4_692      1_521  
  4       31_605      4_051  
  5      156_405     10_794  
  6      120_215     36_468  
  7      268_906     77_972  
  8      443_650     48_714  
  9      358_541     43_348  
 10      307_816     42_707  
 11      292_181     33_415  
 12      290_550     24_693  
 13      199_780     16_644  
 14      158_242     12_939  
 15       76_530     10_411  
 16       56_449      5_385  
 17       22_415      3_340  
 18       20_598      4_879  
 19       10_604      2_404  
 20        3_768      1_514  
 21        2_122       163  
 22         389        50  
 23          74        24  
 24          20         4  
 25           4         0  
~
```

Probabilistic filters

- Inspired by Bloom filter
- Propagate up the tree structure
- Avoid recursion into deeper nesting levels
- Multiple hashing functions
 - Combine several metadata attributes
 - Different phases of filtering
 - Avoid work if not strictly necessary